

## Required R packages and Directories

### Crime Linkage

#### Problem 1: Penalized Regression for Crime Linkage

# Homework #4: Classification

Hyunsuk Ko

Due: Wed Sept 28 | 11:45am

**DS 6030 | Fall 2022 | University of Virginia**

This is an **independent assignment**. Do not discuss or work with classmates.

## Required R packages and Directories

```
data.dir = 'https://mdporter.github.io/DS6030/data/' # data directory
library(R6030)      # functions for SYS-6030
library(tidyverse) # functions for data manipulation
library(glmnet)
library(yardstick)
library(dplyr)
```

## Crime Linkage

Crime linkage attempts to determine if two or more unsolved crimes share a common offender.

*Pairwise* crime linkage is the more simple task of deciding if two crimes share a common offender; it can be considered a binary classification problem. The linkage training data has 8 evidence variables that measure the similarity between a pair of crimes:

- `spatial` is the spatial distance between the crimes
- `temporal` is the fractional time (in days) between the crimes
- `tod` and `dow` are the differences in time of day and day of week between the crimes
- `LOC`, `POA`, and `MOA` are binary with a 1 corresponding to a match (type of property, point of entry, method of entry)
- `TIMERANGE` is the time between the earliest and latest possible times the crime could have occurred (because the victim was away from the house during the crime).
- The response variable indicates if the crimes are linked ( $y = 1$ ) or unlinked ( $y = 0$ ).

These problems use the linkage-train ([https://mdporter.github.io/DS6030/data/linkage\\_train.csv](https://mdporter.github.io/DS6030/data/linkage_train.csv)) and linkage-test ([https://mdporter.github.io/DS6030/data/linkage\\_test.csv](https://mdporter.github.io/DS6030/data/linkage_test.csv)) datasets (click on links for data).

```

train = read.csv('linkage_train.csv')
test = read.csv('linkage_test.csv')

X = glmnet::makeX(
  train = train %>% select(-y),
  test = test
)

X.train = X$x
Y.train = train$y
X.test = X$xtest

```

## Problem 1: Penalized Regression for Crime Linkage

a. Fit a penalized *linear regression* model to predict linkage. Use a lasso, ridge, or elasticnet penalty (your choice).

- Report the value of  $\alpha$  used (if elasticnet)
- Report the value of  $\lambda$  used
- Report the estimated coefficients

```

set.seed(2022)
ridge_cv <- cv.glmnet(X.train, Y.train, alpha = 0, nfolds = 10)
ridge_cv$lambda.min

```

```
#> [1] 0.002327
```

```
coef(ridge_cv, s = "lambda.min")
```

```

#> 9 x 1 sparse Matrix of class "dgCMatrix"
#>                s1
#> (Intercept)  9.263e-02
#> spatial     -2.319e-03
#> temporal    -1.548e-04
#> tod         -2.213e-03
#> dow         -5.795e-03
#> LOC         4.263e-02
#> POA         9.321e-03
#> MOA         7.297e-03
#> TIMERANGE   2.499e-05

```

```
yhat_ridge = predict(ridge_cv, X.test, s = "lambda.min")
```

b. Fit a penalized *logistic regression* model to predict linkage. Use a lasso, ridge, or elasticnet penalty (your choice).

- Report the value of  $\alpha$  used (if elasticnet)
- Report the value of  $\lambda$  used
- Report the estimated coefficients

```
set.seed(2022)
# alpha = 0.5
fit.enet = cv.glmnet(X.train, Y.train, alpha = 0.5, family = "binomial")
fit.enet$lambda.min
```

```
#> [1] 6.912e-05
```

```
coef(fit.enet, s = "lambda.min")
```

```
#> 9 x 1 sparse Matrix of class "dgCMatrix"
#>                s1
#> (Intercept) -0.038162
#> spatial      -0.284236
#> temporal     -0.012414
#> tod          -0.113157
#> dow          -0.240099
#> LOC          1.388922
#> POA          0.441772
#> MOA          0.212107
#> TIMERANGE    0.000615
```

```
yhat_enet = predict(fit.enet, X.test, s = "lambda.min")
```

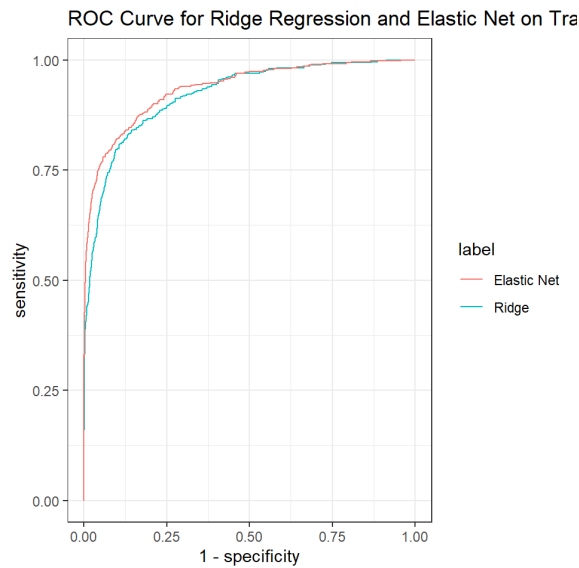
- c. Produce one plot that has the ROC curves, using the *training data*, for both models (from part a and b). Use color and/or linetype to distinguish between models and include a legend.

```
gamma_ridge = predict(ridge_cv, X.train, s = "lambda.min", type = 'link')
gamma_enet = predict(fit.enet, X.train, s = "lambda.min", type = 'link')

ROC_ridge = tibble(truth = factor(Y.train, levels=c(1,0)), gamma = gamma_ridge[,1]) %>%
  yardstick::roc_curve(truth, gamma) %>%
  mutate(label = "Ridge")

ROC_enet = tibble(truth = factor(Y.train, levels=c(1,0)), gamma = gamma_enet[,1]) %>%
  yardstick::roc_curve(truth, gamma) %>%
  mutate(label = "Elastic Net")

ggplot() +
  geom_line(data = ROC_ridge, aes(x = 1-specificity, y = sensitivity, color = label)) +
  geom_line(data = ROC_enet, aes(x = 1-specificity, y = sensitivity, color = label)) +
  labs(title = "ROC Curve for Ridge Regression and Elastic Net on Train Data")
```



d. Recreate the ROC curve from the penalized logistic regression model using repeated hold-out data. The following steps will guide you:

- Fix  $\alpha = .75$
- Run the following steps 25 times:
  - i. Hold out 500 observations
  - ii. Use the remaining observations to estimate  $\lambda$  using 10-fold CV
  - iii. Predict the probability of linkage for the 500 hold-out observations
  - iv. Store the predictions and hold-out labels
- Combine the results and produce the hold-out based ROC curve
- Note: by estimating  $\lambda$  each iteration, we are incorporating the uncertainty present in estimating that tuning parameter.

```

set.seed(2022)

alpha = .75
M = 1#25
K = 10
x = nrow(X.train)
n = 500

holdout_labels = c()
lambdas = c()
linkages = c()
truths = c()
gammas = c()

for (i in 1:M) {
  ind = sample(x, size = n, replace = FALSE)
  holdout = X.train[ind,]
  x_train = X.train[-ind,]
  y_train = Y.train[-ind]

  fit.enet = cv.glmnet(x_train, y_train, alpha = alpha, family = "binomial")
  lambda = fit.enet$lambda.min
  gamma_enet = predict(fit.enet, X.train, s = "lambda.min", type = 'link')

  holdout_labels = append(holdout_labels, i)
  lambdas = append(lambdas, lambda)
  linkages = append(linkages, gamma_enet[,1])
  truths = append(truths, factor(Y.train, levels=c(1,0)))
  gammas = append(gammas, gamma_enet[,1])
}

ROC_tibble = tibble(truth = truths, gamma = gammas)

ROC_enet = ROC_tibble %>%
  yardstick::roc_curve(truth, gamma)

```

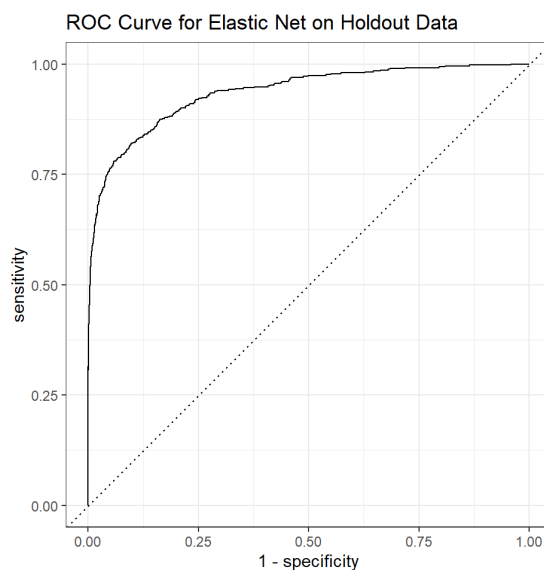
ROC\_enet

<b>.threshold</b> <dbl>	<b>specificity</b> <dbl>	<b>sensitivity</b> <dbl>
-Inf	0.0000000	1.000
-21.265	0.0000000	1.000
-19.876	0.0000339	1.000
-19.791	0.0000678	1.000
-19.643	0.0001017	1.000
-19.444	0.0001356	1.000
-19.321	0.0001695	1.000
-19.295	0.0002034	1.000

.threshold <dbl>	specificity <dbl>	sensitivity <dbl>
-18.974	0.0002373	1.000
-18.879	0.0002712	1.000

1-10 of 10,000 rows      Previous   1   2   3   4   5   6   ...   1000   Next

```
ROC_enet %>%
  ggplot(aes(1-specificity, sensitivity)) + geom_line() +
  geom_abline(lty = 3) +
  coord_equal() +
  labs(title = "ROC Curve for Elastic Net on Holdout Data")
```



```
#table(predicted = ROC_tibble$gamma, truth = ROC_tibble$truth) %>% addmargins()
```

e. Contest Part 1: Predict the estimated *probability* of linkage for the test data (using any model).

- Submit a .csv file (ensure comma separated format) named `lastname_firstname_1.csv` that includes the column named **p** that is your estimated posterior probability. We will use automated evaluation, so the format must be exact.
- You are free to use any tuning parameters
- You are free to use any data transformation or feature engineering
- You will receive credit for a proper submission; the top five scores will receive 2 bonus points.
- Your probabilities will be evaluated with respect to the mean negative Bernoulli log-likelihood (known as the average *log-loss* metric)

$$L = -\frac{1}{M} \sum_{i=1}^m [y_i \log \hat{p}_i + (1 - y_i) \log (1 - \hat{p}_i)]$$

where  $M$  is the number of test observations,  $\hat{p}_i$  is the prediction for the  $i$ th test observation, and  $y_i \in \{0, 1\}$  are the true test set labels.

```
fit.lm = glm(y ~ spatial + temporal + tod + dow + LOC + POA + TIMERANGE, family
= 'binomial', data = train)
```

```
p.hat = predict(fit.lm, test, type = 'response')
```

```
my_phat <- data.frame(p = p.hat)
my_phat
```

	<b>p</b> <dbl>
1	6.287e-02
2	8.713e-03
3	7.362e-02
4	1.257e-05
5	5.127e-04
6	1.043e-02
7	3.609e-02
8	1.852e-02
9	1.216e-02
10	2.843e-05
1-10 of 10,000 rows	Previous 1 2 3 4 5 6 ... 1000 Next

```
write_csv(my_phat, 'ko_hyunsuk_1.csv')
```

f. Contest Part 2: Predict the linkages for the test data (using any model).

- Submit a .csv file (ensure comma separated format) named `lastname_firstname_2.csv` that includes the column named **linkage** that takes the value of 1 for linkages and 0 for unlinked pairs. We will use automated evaluation, so the format must be exact.
- You are free to use any tuning parameters.
- You are free to use any data transformation or feature engineering.
- Your labels will be evaluated based on total cost, where cost is equal to  $1 \cdot \text{FP} + 8 \cdot \text{FN}$ . This implies that False Negatives (FN) are 8 times as costly as False Positives (FP)
- You will receive credit for a proper submission; the top five scores will receive 2 bonus points. Note: you only will get bonus credit for one of the two contests.

```
G.hat = ifelse(p.hat >= .1, 1, 0)
```

```
my_ghat <- data.frame(linkage = G.hat)
my_ghat
```

**linkage**  
<dbl>

	linkage
	<dbl>
1	0
2	0
3	0
4	0
5	0
6	0
7	0
8	0
9	0
10	0
1-10 of 10,000 rows	Previous 1 2 3 4 5 6 ... 1000 Next

```
write_csv(my_ghat, 'ko_hyunsuk_2.csv')
```